

# One Pixel Attack for Fooling Deep Neural Networks

Jiawei Su\*, Danilo Vasconcellos Vargas\* and Kouichi Sakurai

**Abstract**—Recent research has revealed that the output of Deep Neural Networks (DNN) can be easily altered by adding relatively small perturbations to the input vector. In this paper, we analyze an attack in an extremely limited scenario where only one pixel can be modified. For that we propose a novel method for generating one-pixel adversarial perturbations based on differential evolution (DE). It requires less adversarial information (a black-box attack) and can fool more types of networks due to the inherent features of DE. The results show that 68.36% of the natural images in CIFAR-10 test dataset and 41.22% of the ImageNet (ILSVRC 2012) validation images can be perturbed to at least one target class by modifying just one pixel with 73.22% and 5.52% confidence on average. Thus, the proposed attack explores a different take on adversarial machine learning in an extreme limited scenario, showing that current DNNs are also vulnerable to such low dimension attacks. Besides, we also illustrate an important application of DE (or broadly speaking, evolutionary computation) in the domain of adversarial machine learning: creating tools that can effectively generate low-cost adversarial attacks against neural networks for evaluating robustness.

**Index Terms**—Differential Evolution, Convolutional Neural Network, Information Security, Image Recognition.

## I. INTRODUCTION

IN the domain of image recognition, DNN-based approach has outperformed traditional image processing techniques, achieving even human-competitive results [25]. However, several studies have revealed that artificial perturbations on natural images can easily make DNN misclassify and accordingly proposed effective algorithms for generating such samples called “adversarial images” [7][11][18][24]. A common idea for creating adversarial images is adding a tiny amount of well-tuned additive perturbation, which is expected to be imperceptible to human eyes, to a correctly classified natural image. Such modification can cause the classifier to label the modified image as a completely different class. Unfortunately, most of the previous attacks did not consider extremely limited scenarios for adversarial attacks, namely the modifications might be excessive (i.e., the amount of modified pixels is fairly large) such that it may be perceptible to human eyes (see Figure 3 for an example). Additionally, investigating adversarial images created under extremely limited scenarios might give new insights about the geometrical characteristics and overall behavior of DNN’s model in high dimensional space [9]. For

Authors are with the Graduate School/Faculty of Information Science and Electrical Engineering, Kyushu University, Japan. The third author is also affiliated to Advanced Telecommunications Research Institute International (ATR).

The code is available on: <https://github.com/Carina02/One-Pixel-Attack>

\*Both authors have equal contribution.




Fig. 1. One-pixel attacks created with the proposed algorithm that successfully fooled three types of DNNs trained on CIFAR-10 dataset: The All convolutional network (AllConv), Network in network (NiN) and VGG. The original class labels are in black color while the target class labels and the corresponding confidence are given below.

example, the characteristics of adversarial images close to the decision boundaries can help describing the boundaries’ shape.

In this paper, by perturbing only one pixel with differential evolution, we propose a black-box DNN attack in a scenario where the only information available is the probability labels (Figure 1 and 2). Our proposal has mainly the following advantages compared to previous works:

- **Effectiveness** - On CIFAR-10 dataset, being able to launch non-targeted attacks by only modifying one pixel

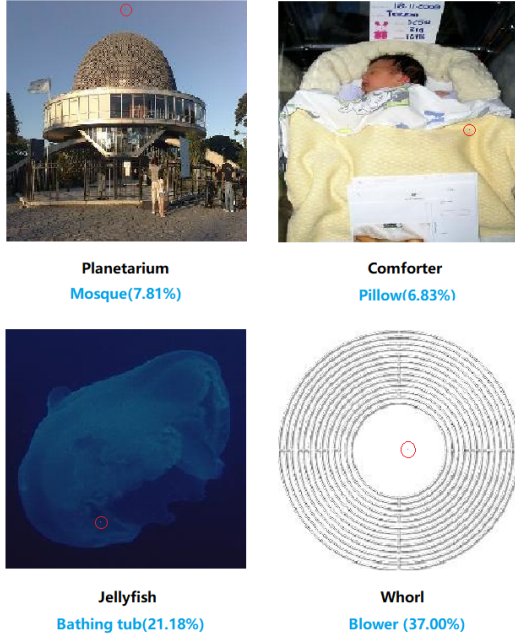


Fig. 2. One-pixel attacks on ImageNet dataset where the modified pixels are highlighted with red circles. The original class labels are in black color while the target class labels and their corresponding confidence are given below.

on three common deep neural network structures with 68.71%, 72.85% and 63.53% success rates. We additionally find that each natural image can be perturbed to 1.9, 2.3 and 1.7 other classes. While on ImageNet dataset, non-targeted attacking the BVL AlexNet model also by changing one pixel shows that 41.22% of the validation images can be attacked.

- **Semi-Black-Box Attack** - Requires only black-box feedback (probability labels) but no inner information of target DNNs such as gradients and network structures. Our method is also simpler than existing approaches since it does not abstract the problem of searching perturbation to any explicit target functions but directly focus on increasing the probability label values of the target classes.
- **Flexibility** - Can attack more types of DNNs (e.g., networks that are not differentiable or when the gradient calculation is difficult).

Regarding the extremely limited one-pixel attack scenario, there are several main reasons why we consider it:

- **Analyze the Vicinity of Natural Images** - Geometrically, several previous works have analyzed the vicinity of natural images by limiting the length of perturbation vector. For example, the universal perturbation adds small value to each pixel such that it searches the adversarial images in a sphere region around the natural image [14]. On the other side, the proposed few-pixel perturbations can be regarded as cutting the input space using very low-dimensional slices, which is a different way of exploring the features of high dimensional DNN input space. Among them, one-pixel attack is an extreme case of several-pixel attack. Theoretically, it can give geometrical




Fig. 3. An illustration of the adversarial images generated by using Jacobian saliency-map approach [18]. The perturbation is conducted on about 4% of the total pixels and can be obvious to human eyes. Since the adversarial pixel perturbation has become a common way of generating adversarial images, such abnormal “noise” might be recognized with expertise.

insight to the understanding of CNN input space, in contrast to another extreme case: universal adversarial perturbation [14] that modifies every pixel.

- **A Measure of Perceptiveness** - The attack can be effective for hiding adversarial modification in practice. To the best of our knowledge, none of the previous works can guarantee that the perturbation made can be completely imperceptible. A direct way of mitigating this problem is to limit the amount of modifications to as few as possible. Specifically, instead of theoretically proposing additional constraints or considering more complex cost functions for conducting perturbation, we propose an empirical solution by limiting the number of pixels that can be modified. In other words, we use the number of pixels as units instead of length of perturbation vector to measure the perturbation strength and consider the worst case which is one-pixel modification, as well as two other scenarios (i.e. 3 and 5 pixels) for comparison.

## II. RELATED WORKS

The security problem of DNN has become a critical topic [1][2]. C. Szegedy et al. first revealed the sensitivity to well-tuned artificial perturbation [24] which can be crafted by several gradient-based algorithms using back-propagation for obtaining gradient information [11][24]. Specifically, I.J. Goodfellow et al. proposed “fast gradient sign” algorithm for calculating effective perturbation based on a hypothesis in which the linearity and high-dimensions of inputs are the main reason that a broad class of networks are sensitive to small perturbation [11]. S.M. Moosavi-Dezfooli et al. proposed a greedy perturbation searching method by assuming the linearity of DNN decision boundaries [7]. In addition, N.

Papernot et al. utilize Jacobian matrix to build “Adversarial Saliency Map” which indicates the effectiveness of conducting a fixed length perturbation through the direction of each axis [18][20]. Except adversarial perturbation, there are other ways of creating adversarial images to make the DNN misclassify, such as artificial image [16] and rotation [36]. Besides, adversarial perturbation can be also possible in other domains such as speech recognition [33], natural language processing [34] and malware classification [35].

A number of detection and defense methods have been also proposed to mitigate the vulnerability induced by adversarial perturbation [39]. For instance, network distillation which was originally proposed for squeezing information of an network to a smaller one is found to be able to reduce the network sensitivity enhancing the robustness of the neural network [40]. Adversarial training [41] is proposed for adding adversarial images to the training data such that the robustness against known adversarial images can be improved. On the other side, some image processing methods are proved to be effective for detecting adversarial images. For example, B.Liang et al. show that noise reduction methods such as scalar quantization and spatial smoothing filter can be selectively utilized for mitigating the influence of adversarial perturbation. By comparing the label of an image before and after the transformation the perturbation can be detected [43]. The method works well on detecting adversarial images with both low and high entropy. Similarly, W. Xu et al. show that squeezing color bits and local/non-local spatial smoothing can have high success rate on distinguishing adversarial images [42]. However, recent studies show that many of these defense and detection methods can be effectively evaded by conducting little modification on the original attacks [45], [46], [60].

Several black-box attacks that require no internal knowledge about the target systems such as gradients, have also been proposed [5][17][15]. In particular, to the best of our knowledge, the only work before ours that ever mentioned using one-pixel modification to change class labels is carried out by N. Narodytska et al[15]. However, differently from our work, they only utilized it as a starting point to derive a further semi black-box attack which needs to modify more pixels (e.g., about 30 pixels out of 1024) without considering the scenario of one-pixel attack. In addition, they have neither measured systematically the effectiveness of the attack nor obtained quantitative results for evaluation. An analysis of the one-pixel attack’s geometrical features as well as further discussion about its implications are also lacking.

There have been many efforts to understand DNN by visualizing the activation of network nodes [28] [29][30] while the geometrical characteristics of DNN boundary have gained less attraction due to the difficulty of understanding high-dimensional space. However, the robustness evaluation of DNN with respect to adversarial perturbation might shed light in this complex problem [9]. For example, both natural and random images are found to be vulnerable to adversarial perturbation. Assuming these images are evenly distributed, it suggests that most data points in the input space are gathered near to the boundaries [9]. In addition, A. Fawzi et al. revealed more clues by conducting a curvature analysis. Their conclu-

sion is that the region along most directions around natural images are flat with only few directions where the space is curved and the images are sensitive to perturbation[10]. Interestingly, universal perturbations (i.e. a perturbation that when added to any natural image can generate adversarial images with high effectiveness) were shown possible and to achieve a high effectiveness when compared to random perturbation. This indicates that the diversity of boundaries might be low while the boundaries’ shapes near different data points are similar [14].

### III. METHODOLOGY

#### A. Problem Description

Generating adversarial images can be formalized as an optimization problem with constraints. We assume an input image can be represented by a vector in which each scalar element represents one pixel. Let  $f$  be the target image classifier which receives  $n$ -dimensional inputs,  $\mathbf{x} = (x_1, \dots, x_n)$  be the original natural image correctly classified as class  $t$ . The probability of  $\mathbf{x}$  belonging to the class  $t$  is therefore  $f_t(\mathbf{x})$ . The vector  $e(\mathbf{x}) = (e_1, \dots, e_n)$  is an additive adversarial perturbation according to  $\mathbf{x}$ , the target class  $adv$  and the limitation of maximum modification  $L$ . Note that  $L$  is always measured by the length of vector  $e(\mathbf{x})$ . The goal of adversaries in the case of targeted attacks is to find the optimized solution  $e(\mathbf{x})^*$  for the following question:

$$\begin{aligned} & \underset{e(\mathbf{x})^*}{\text{maximize}} && f_{adv}(\mathbf{x} + e(\mathbf{x})) \\ & \text{subject to} && \|e(\mathbf{x})\| \leq L \end{aligned}$$

The problem involves finding two values: (a) which dimensions that need to be perturbed and (b) the corresponding strength of the modification for each dimension. In our approach, the equation is slightly different:

$$\begin{aligned} & \underset{e(\mathbf{x})^*}{\text{maximize}} && f_{adv}(\mathbf{x} + e(\mathbf{x})) \\ & \text{subject to} && \|e(\mathbf{x})\|_0 \leq d, \end{aligned}$$

where  $d$  is a small number. In the case of one-pixel attack  $d = 1$ . Previous works commonly modify a part of all dimensions while in our approach only  $d$  dimensions are modified with the other dimensions of  $e(\mathbf{x})$  left to zeros.

The one-pixel modification can be seen as perturbing the data point along a direction parallel to the axis of one of the  $n$  dimensions. Similarly, the 3(5)-pixel modification moves the data points within 3(5)-dimensional cubes. Overall, few-pixel attack conducts perturbations on the low-dimensional slices of input space. In fact, one-pixel perturbation allows the modification of an image towards a chosen direction out of  $n$  possible directions with arbitrary strength. This is illustrated in Figure 4 for the case when  $n = 3$ .

Thus, usual adversarial images are constructed by perturbing all pixels with an overall constraint on the strength of accumulated modification[8][14] while the few-pixel attack considered in this paper is the opposite which specifically focus on few pixels but does not limit the strength of modification.




Fig. 4. An illustration of using one and two-pixel perturbation attack in a 3-dimensional input space (i.e. the image has three pixels). The green point (sphere) denotes a natural image. In the case of one-pixel perturbation, the search space is the three perpendicular lines that intersect at point of natural image, which are denoted by red and black stripes. For two-pixel perturbation, the search space is the three blue (shaded) two-dimensional planes. In summary, one and two-pixel attacks search the perturbation on respectively one and two dimensional slices of the original three dimensional input space.

### B. Differential Evolution

Differential evolution (DE) is a population based optimization algorithm for solving complex multi-modal optimization problems [23], [6]. DE belongs to the general class of evolutionary algorithms (EA). Moreover, it has mechanisms in the population selection phase that keep the diversity such that in practice it is expected to efficiently find higher quality solutions than gradient-based solutions or even other kinds of EAs [4]. In specific, during each iteration another set of candidate solutions (children) is generated according to the current population (parents). Then the children are compared with their corresponding parents, surviving if they are more fitted (possess higher fitness value) than their parents. In such a way, only comparing the parent and his child, the goal of keeping diversity and improving fitness values can be simultaneously achieved.

DE does not use the gradient information for optimizing and therefore does not require the objective function to be differentiable or previously known. Thus, it can be utilized on a wider range of optimization problems compared to gradient based methods (e.g., non-differentiable, dynamic, noisy, among others). The use of DE for generating adversarial images have the following main advantages:

- **Higher probability of Finding Global Optima** - DE is a meta-heuristic which is relatively less subject to local minima than gradient descent or greedy search algorithms (this is in part due to diversity keeping mechanisms and the use of a set of candidate solutions). Moreover, the problem considered in this article has a strict constraint (only one pixel can be modified) making it relatively harder.
- **Require Less Information from Target System** - DE does not require the optimization problem to be differentiable as is required by classical optimization methods such as gradient descent and quasi-newton methods. This

is critical in the case of generating adversarial images since 1) There are networks that are not differentiable, for instance [26]. 2) Calculating gradient requires much more information about the target system which can be hardly realistic in many cases.

- **Simplicity** - The approach proposed here is independent of the classifier used. For the attack to take place it is sufficient to know the probability labels.

There are many DE variations/improvements such as self-adaptive [3], multi-objective [27], among others. The current work can be further improved by taking these variations/improvements into account.

### C. Method and Settings

We encode the perturbation into an array (candidate solution) which is optimized (evolved) by differential evolution. One candidate solution contains a fixed number of perturbations and each perturbation is a tuple holding five elements: x-y coordinates and RGB value of the perturbation. One perturbation modifies one pixel. The initial number of candidate solutions (population) is 400 and at each iteration another 400 candidate solutions (children) will be produced by using the usual DE formula:

$$x_i(g+1) = x_{r1}(g) + F(x_{r2}(g) + x_{r3}(g)), \\ r1 \neq r2 \neq r3,$$

where  $x_i$  is an element of the candidate solution,  $r1, r2, r3$  are random numbers,  $F$  is the scale parameter set to be 0.5,  $g$  is the current index of generation. Once generated, each candidate solution compete with their corresponding parents according to the index of the population and the winner survive for next iteration. The maximum number of iteration is set to 100 and early-stop criteria will be triggered when the probability label of target class exceeds 50% in the case of targeted attacks on CIFAR-10, and when the label of true class is lower than 5% in the case of non-targeted attacks on ImageNet. Then the label of true class is compared with the highest non-true class to evaluate if the attack succeeded. The initial population is initialized by using uniform distributions  $U(1, 32)$  for CIFAR-10 images and  $U(1, 227)$  for ImageNet images, for generating x-y coordinate (e.g., the image has a size of 32X32 in CIFAR-10 and for ImageNet we unify the original images with various resolutions to 227X227) and Gaussian distributions  $N(\mu=128, \sigma=127)$  for RGB values. The fitness function is simply the probabilistic label of the target class in the case of CIFAR-10 and the label of true class in the case of ImageNet. The crossover is not included in our scheme.

## IV. EVALUATION AND RESULTS

The evaluation of the proposed attack method is based on CIFAR-10 and ImageNet dataset. We introduce several metrics to measure the effectiveness of the attacks:

- **Success Rate** - In the case of non-targeted attacks, it is defined as the percentage of adversarial images that were

successfully classified by the target system as an arbitrary target class. In the case of targeted attack, it is defined as the probability of perturbing a natural image to a specific target class.

- **Adversarial Probability Labels(Confidence)** - Accumulates the values of probability label of the target class for each successful perturbation, then divided by the total number of successful perturbations. The measure indicates the average confidence given by the target system when mis-classifying adversarial images.
- **Number of Target Classes** - Counts the number of natural images that successfully perturb to a certain number (i.e. from 0 to 9) of target classes. In particular, by counting the number of images that can not be perturbed to any other classes, the effectiveness of non-targeted attack can be evaluated.
- **Number of Original-Target Class Pairs** - Counts the number of times each original-destination class pair was attacked.

### A. CIFAR-10

We train 3 types of common networks: All convolution network [22], Network in Network[13] and VGG16 network[21] as target image classifiers on cifar-10 dataset [12]. The structures of the networks are described in Table 1, 2 and 3. The network setting were kept as similar as possible to the original with a few modifications in order to get the highest classification accuracy. Both the scenarios of targeted and non-targeted attacks are considered. For each of the attacks on the three types of neural networks 500 natural images are randomly selected from the cifar-10 test dataset to conduct the attack. In addition, an experiment is conducted on the all convolution network [22] by generating 500 adversarial images with three and five pixel-modification. The objective is to compare one-pixel attack with three and five pixel attacks. For each natural image, nine target attacks are launched trying to perturb it to the other 9 target classes. Note that we actually only launch targeted attacks and the effectiveness of non-targeted attack is evaluated based on targeted attack results. That is, if an image can be perturbed to at least one target class out of total 9 classes, the non-targeted attack on this image succeeds. Overall, it leads to the total of 36000 adversarial images created. To evaluate the effectiveness of the attacks, some established measures from the literature are used as well as some new kinds of measures are introduced:

### B. ImageNet

For ImageNet we applied a non-targeted attack with the same DE parameter settings used on the CIFAR-10 dataset, although ImageNet has a search space 50 times larger than CIFAR-10. Note that we actually launch the non-targeted attack for ImageNet by using a fitness function that aims to decrease the probability label of the true class. Different from CIFAR-10, whose effectiveness of non-targeted attack is calculated based on the targeted attack results carried out by using a fitness function for increasing the probability of target classes. Given the time constraints, we conduct the experiment

```
conv2d layer(kernel=3, stride = 1, depth=96)
conv2d layer(kernel=3, stride = 1, depth=96)
conv2d layer(kernel=3, stride = 2, depth=96)
conv2d layer(kernel=3, stride = 1, depth=192)
conv2d layer(kernel=3, stride = 1, depth=192)
dropout(0.3)
conv2d layer(kernel=3, stride = 2, depth=192)
conv2d layer(kernel=3, stride = 2, depth=192)
conv2d layer(kernel=1, stride = 1, depth=192)
conv2d layer(kernel=1, stride = 1, depth=10)
average pooling layer(kernel=6, stride=1)
flatten layer
softmax classifier
```

TABLE I  
ALL CONVOLUTION NETWORK

```
conv2d layer(kernel=5, stride = 1, depth=192)
conv2d layer(kernel=1, stride = 1, depth=160)
conv2d layer(kernel=1, stride = 1, depth=96)
max pooling layer(kernel=3, stride=2)
dropout(0.5)
conv2d layer(kernel=5, stride = 1, depth=192)
conv2d layer(kernel=5, stride = 1, depth=192)
conv2d layer(kernel=5, stride = 1, depth=192)
average pooling layer(kernel=3, stride=2)
dropout(0.5)
conv2d layer(kernel=3, stride = 1, depth=192)
conv2d layer(kernel=1, stride = 1, depth=192)
conv2d layer(kernel=1, stride = 1, depth=10)
flatten layer
softmax classifier
```

TABLE II  
NETWORK IN NETWORK

```
conv2d layer(kernel=3, stride = 1, depth=64)
conv2d layer(kernel=3, stride = 1, depth=64)
max pooling layer(kernel=2, stride=2)
conv2d layer(kernel=3, stride = 1, depth=128)
conv2d layer(kernel=3, stride = 1, depth=128)
max pooling layer(kernel=2, stride=2)
conv2d layer(kernel=3, stride = 1, depth=256)
conv2d layer(kernel=3, stride = 1, depth=256)
conv2d layer(kernel=3, stride = 1, depth=256)
max pooling layer(kernel=2, stride=2)
conv2d layer(kernel=3, stride = 1, depth=512)
conv2d layer(kernel=3, stride = 1, depth=512)
conv2d layer(kernel=3, stride = 1, depth=512)
max pooling layer(kernel=2, stride=2)
conv2d layer(kernel=3, stride = 1, depth=512)
conv2d layer(kernel=3, stride = 1, depth=512)
conv2d layer(kernel=3, stride = 1, depth=512)
max pooling layer(kernel=2, stride=2)
flatten layer
fully connected(size=2048)
fully connected(size=2048)
softmax classifier
```

TABLE III  
VGG16 NETWORK

without proportionally increasing the number of evaluations, i.e. we keep the same number of evaluations. Our tests are run over the BVLC AlexNet using 600 images from ILSVRC 2012 validation set selected randomly for the attack. For ImageNet we only conduct one pixel attack because we are want to verify if such a tiny modification can fool images with larger size and if it is computationally tractable to conduct such attacks.

### C. Results

The success rates and adversarial probability labels for one-pixel perturbations on three CIFAR-10 networks and BVLC



	AllConv	NiN	VGG16	BVLC
OriginAcc	86.94%	87.70%	85.45%	57.26%
Targeted	19.82%	23.15%	17.09%	–
Non-targeted	68.71%	72.85%	63.53%	41.23%
Confidence	79.40%	75.02%	65.25%	5.53%

TABLE IV

RESULTS OF CONDUCTING ONE-PIXEL ATTACK ON FOUR DIFFERENT TYPES OF NETWORKS: ALL CONVOLUTIONAL NETWORK(ALLCONV), NETWORK IN NETWORK(NiN), VGG16 AND BVLC ALEXNET. THE ORIGINALACC IS THE ACCURACY ON THE NATURAL TEST DATASETS. TARGETED/NON-TARGETED INDICATE THE ACCURACY OF CONDUCTING TARGETED/NON-TARGETED ATTACKS. CONFIDENCE IS THE AVERAGE PROBABILITY OF TARGET CLASSES.

network are shown in Table 4 and the three and five-pixel perturbations on CIFAR-10 is shown in Table 5. The number of target classes is shown by Figure 5. The number of original-target class pairs is shown by the heat-maps of Figure 6 and 7. In addition to the number of original-target class pairs, the total number of times each class had an attack which either originated or targeted it is shown in Figure 8. Since only non-targeted attacks are launched on ImageNet, the “Number of target classes” and “Number of original-target class pairs” metrics are not included in the ImageNet results.

1) *Success Rate and Adversarial Probability Labels (Targeted Attack Results)*: On CIFAR-10, the success rates of one-pixel attacks on three types of networks show the generalized effectiveness of the proposed attack through different network structures. On average, each image can be perturbed to about two target classes for each network. In addition, by increasing the number of pixels that can be modified to three and five, the number of target classes that can be reached increases significantly. By dividing the adversarial probability labels by the success rates, the confidence values (i.e. probability labels of target classes) are obtained which are 79.39%, 79.17% and 77.09% respectively to one, three and five-pixel attacks.

On ImageNet, the results show that the one pixel attack generalizes well to large size images and fool the corresponding neural networks. In particular, there is 41.23% chance that an arbitrary ImageNet validation image can be perturbed to a target class with 5.53% confidence. Note that the ImageNet results are done with the same settings as CIFAR-10 while the resolution of images we use for the ImageNet test is 227x227, which is 50 times larger than CIFAR-10 (32x32). Notice that in each successful attack the probability label of the target class is the highest. Therefore, the confidence of 5.53% is relatively low but tell us that the other remaining 999 classes are even lower to an almost uniform soft label distribution. Thus, the one-pixel attack can break the confidence of BVLC AlexNet to a nearly uniform soft label distribution. The low confidence is caused by the fact that we utilized a non-targeted evaluation that only focuses on decreasing the probability of the true class. Other fitness functions should give different results.

2) *Number of Target Classes (Non-targeted Attack Results)*: Regarding the results shown in Figure 5, we find that with only one-pixel modification a fair amount of natural images can be perturbed to two, three and four target classes. By increasing the number of pixels modified, perturbation to more target

	3 pixels	5 pixels
Success rate(tar)	40.57%	44.00%
Success rate(non-tar)	86.53%	86.34%
Rate/Labels	79.17%	77.09%

TABLE V

RESULTS OF CONDUCTING THREE-PIXEL ATTACK ON ALLCONV NETWORKS AND FIVE-PIXEL ATTACK ON NETWORK IN NETWORK.

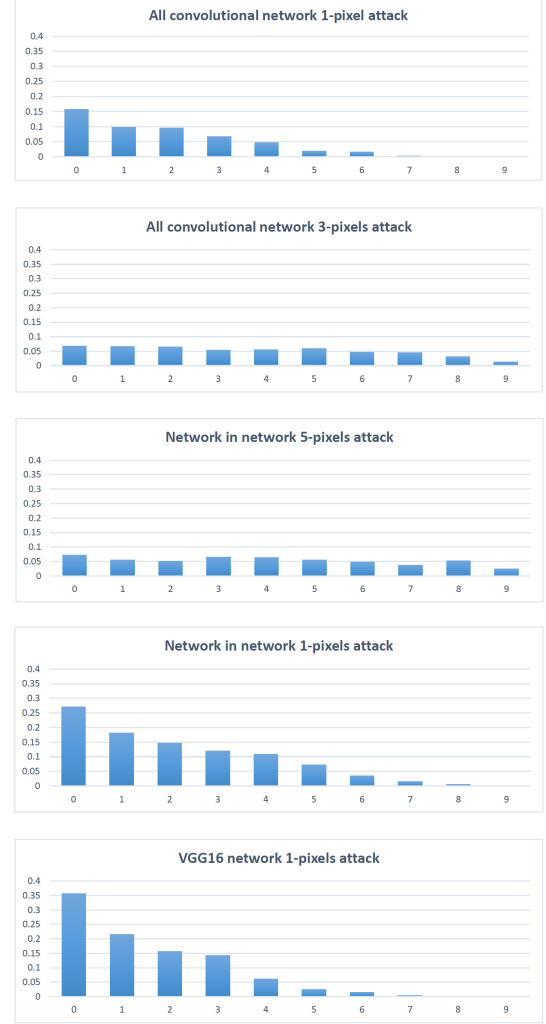


Fig. 5. The graphs shows the percentage of natural images that were successfully perturbed to a certain number (from 0 to 9) of target classes by using one, three or five-pixel perturbation. The vertical axis shows the percentage of images that can be perturbed while the horizontal axis indicates the number of target classes.

classes becomes highly probable. In the case of non-targeted one-pixel attack, the VGG16 network got a slightly higher robustness against the proposed attack. This suggests that all three types of networks (AllConv network, NiN and VGG16) are vulnerable to this type of attack.

The results of attacks are competitive with previous non-targeted attack methods which need much more distortions (Table 6). It shows that using one dimensional perturbation vectors is enough to find the corresponding adversarial images for most of the natural images. In fact, by increasing the number of pixels up to five, a considerable number of images

Method	Success rate	Confidence	Number of pixels	Network
Our method	72.85%	75%	1 (0.098%)	NiN
Our method	63.53%	65%	1 (0.098%)	VGG
Our method	41.23%	6%	1 (0.002%)	BVLG
LSA[15]	97.89%	72%	33 (3.24%)	NiN
LSA[15]	97.98%	77%	30 (2.99%)	VGG
FGSM[11]	93.67%	93%	1024 (100%)	NiN
FGSM[11]	90.93%	90%	1024 (100%)	VGG

TABLE VI

COMPARISON OF NON-TARGETED ATTACK EFFECTIVENESS BETWEEN THE PROPOSED METHOD AND TWO PREVIOUS WORKS. THIS SUGGESTS THAT ONE PIXEL IS ENOUGH TO CREATE ADVERSARIAL IMAGES FROM MOST OF THE NATURAL IMAGES.

can be simultaneously perturbed to eight target classes. In some rare cases, an image can go to all other target classes with one-pixel modification, which is illustrated in Figure 9.

3) *Original-Target Class Pairs*: Some specific original-target class pairs are much more vulnerable than others (Figure 6 and 7). For example, images of cat (class 3) can be much more easily perturbed to dog (class 5) but can hardly reach the automobile (class 1). This indicates that the vulnerable target classes (directions) are shared by different data points that belong to the same class. Moreover, in the case of one-pixel attack, some classes are more robust than others since their data points can be relatively hard to perturb to other classes. Among these data points, there are points that can not be perturbed to any other classes. This indicates that the labels of these points rarely change when going across the input space through  $n$  directions perpendicular to the axes. Therefore, the corresponding original classes are kept robust along these directions. However, it can be seen that such robustness can rather easily be broken by merely increasing the dimensions of perturbation from one to three and five because both success rates and number of target classes that can be reached increase when conducting higher-dimensional perturbations.

Additionally, it can also be seen that each heat-map matrix is approximately symmetric, indicating that each class has similar number of adversarial images which were crafted from these classes as well as to these classes (Figure 8). Having said that, there are some exceptions for example the class 8 (ship) when attacking NiN, the class 4 (deer) when attacking AllConv networks with one pixel, among others. In the ship class when attacking NiN networks, for example, it is relatively easy to craft adversarial images from them while it is relatively hard to craft adversarial images to them. Such unbalance is intriguing since it indicates the ship class is similar to most of the other classes like truck and airplane but not vice-versa. This might be due to (a) boundary shape and (b) how close are natural images to the boundary. In other words, if the boundary shape is wide enough it is possible to have natural images far away from the boundary such that it is hard to craft adversarial images from it. On the contrary, if the boundary shape is mostly long and thin with natural images close to the border, it is easy to craft adversarial images from them but hard to




Fig. 6. Heat-maps of the number of times a successful attack is present with the corresponding original-target class pair in one, three and five-pixel attack cases. Red (vertical) and (horizontal) blue indices indicate respectively the original and target classes. The number from 0 to 9 indicates respectively the following classes: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.

craft adversarial images to them.

In practice, such classes which are easy to craft adversarial images from may be exploited by malicious users which may make the whole system vulnerable. In the case here, however, the exceptions are not shared between the networks, revealing that whatever is causing the phenomenon is not shared. Therefore, for the current systems under the given attacks, such a vulnerability seems hard to be exploited.

4) *Time complexity and average distortion*: To evaluate the time complexity we use the number of evaluations which is a common metric in optimization. In the DE case the number of evaluations is equal to the population size multiplied by the number of generations. We also calculate the average distortion




Fig. 7. Heat-maps for one-pixel attack on Network in network and VGG.

	AllConv	NiN	VGG16	BVLC
AvgEvaluation	16000	12400	20000	25600
AvgDistortion	100	114	115	101

TABLE VII

COST OF CONDUCTING ONE-PIXEL ATTACK ON FOUR DIFFERENT TYPES OF NETWORKS. AVGEVALUATION IS THE AVERAGE NUMBER OF EVALUATIONS TO PRODUCE ADVERSARIAL IMAGES. AVGDISTORTION IS THE REQUIRED AVERAGE DISTORTION IN ONE-CHANNEL OF A SINGLE PIXEL TO PRODUCE ADVERSARIAL IMAGES.

on the single pixel attacked by taking the average modification on the three color channels, which is a more straight forward and explicit measure of modification strength. We did not use the  $L_p$  norm due to its limited effectiveness of measuring perceptiveness [38]. The results of two metrics are shown in Table 7.

5) *Comparing with Random One-Pixel Attack:* We compare the proposed method with the random attack to evaluate if DE is truly helpful for conducting one-pixel non-targeted attack on CIFAR-10 dataset, which is shown in Table 8.

Specifically, for each natural image, the random search repeats 100 times, each time randomly modifies one random pixel of the image with random RGB value to attempt to change its label. The confidence of the attack with respect to one image is set to be the highest probability target class label of 100 attacks.

In this experiment, we use the same number of evaluations (80000) for both DE and random search. According to the comparison, the DE is superior to the random attack regarding attack accuracy, especially in the case of VGG16 network.




Fig. 8. Number of successful attacks (vertical axis) for a specific class acting as the original (black) and target (gray) class. The horizontal axis indicates the index of each class which is the same as Figure 7.

Specifically, DE is 19.01%, 31.13% and 47.96% more efficient than random search respectively for All Convolutional Network, Network in Network and VGG16. Even with a less efficient result, random search is shown to find 49.70% and 41.72% of the time for respectively All Convolutional Network and Network in Network, therefore the vulnerable pixels that can change the image label significantly are quite common. That seems not to be the case for VGG though in which random search achieves only 15.57%. DE has a similar





Fig. 9. A natural image of the dog class that can be perturbed to all other nine classes. The attack is conducted over the AllConv network using the proposed one pixel attack. The table in the bottom shows the class labels output by the target DNN, all with approximately 100% confidence. This curious result further emphasize the difference and limitations of current methods when compared to human recognition.

	AllConv	NiN	VGG16
DE success rate	68.71%	72.85%	63.53%
Confidence	79.40%	75.02%	65.25%
Random Search success rate	49.70%	41.72%	15.57%
Confidence	87.73%	75.83%	59.90%

TABLE VIII  
A COMPARISON OF ATTACK RATE AND CONFIDENCE BETWEEN DE ONE-PIXEL ATTACK AND RANDOM ONE-PIXEL ATTACK(NON-TARGETED) ON CIFAR-10 DATASET.

accuracy in all of them showing also a better robustness.

6) *Change in fitness values*: We run an experiment over different networks to examine how the fitness changes during evolution. The 30 curves come from 30 random images attacked by the proposed one-pixel attack (Figure 10). The fitness values are, as previously described, set to be the probability label of the true class for each image. The goal of the attack is to minimize this fitness value. According to the results, it can be seen that the fitness values can occasionally drop abruptly between two generations while in other cases they decrease smoothly. Moreover, the average fitness value decreases monotonically with the number of generations, showing that the evolution works as expected. We also find that BVLC network is harder to fool due to the smaller decrease in fitness values.

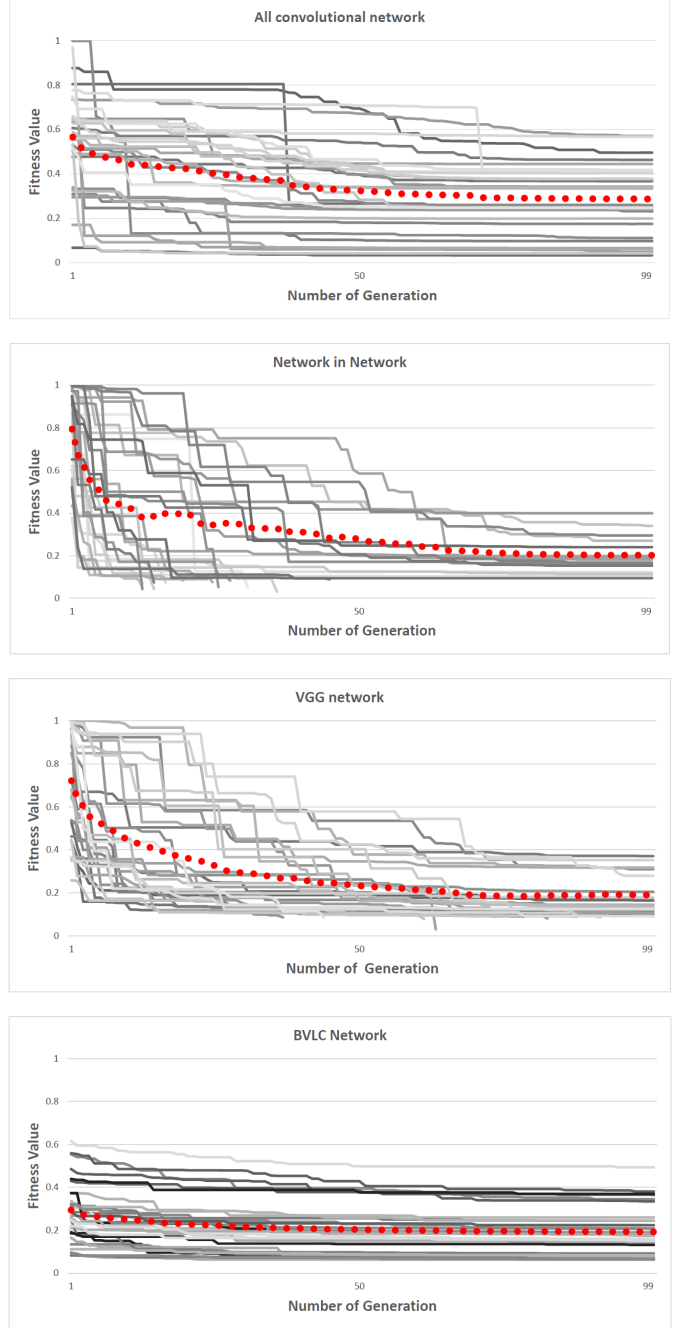


Fig. 10. The change of fitness values during 100 generations of evolution of 30 random images (non-targeted) attacked by the proposed method among different network structures. The average values are highlighted by red dotted lines.

## V. DISCUSSION

### A. Adversarial Perturbation

Previous results have shown that many data points might be located near to the decision boundaries [9]. For the analysis the data points were moved small steps in the input space while quantitatively analyzing the frequency of change in the class labels. In this paper, we showed that it is also possible to move the data points along few dimension to find points where the class labels change. Our results also suggest that the

assumption made by I. J. Goodfellow et al. that small additive perturbation on the values of many dimensions will accumulate and cause huge change to the output [11], might not be necessary for explaining why natural images are sensitive to small perturbation. Since we only changed one pixel to successfully perturb a considerable number of images.

According to the experimental results, the vulnerability of CNN exploited by the proposed one pixel attack is generalized through different network structures as well as different image sizes. In addition, the results shown here mimics an attacker and therefore uses a low number of DE iterations with a relatively small set of initial candidate solutions. Therefore, the perturbation success rates should improve further by having either more iterations or a bigger set of initial candidate solutions. Implementing more advanced algorithms such as Co-variance Matrix Adaptation Evolution Strategy [32] instead of DE might also achieve the same improvement. Additionally, the proposed algorithm and the widely vulnerable images (i.e. natural images that can be used to craft adversarial images to most of the other classes) collected might be useful for generating better artificial adversarial images in order to augment the training data set. This aids the development of more robust models[19] which is left as future work.

### B. Robustness of One-pixel Attack

Some recently proposed detection methods have shown high accuracy of detecting adversarial perturbation. For example, B.Liang et al. utilize noise reduction to effectively detect both high and low-entropy images (e.g., bigger images give high entropy values) [43]. In addition, W. Xu et al. show that squeezing color bits and local/non-local spatial smoothing can simultaneously detect  $L_0$ ,  $L_2$  and  $L_\infty$  attacks [42]. As the trade-off of being a low-cost, easy-implemented  $L_0$  attack, we do not expect one pixel attack can achieve significantly better robustness against such detection methods compared to other  $L_0$  attacks such as [32].

However, such detection schemes add another layer of pre-processing which increases the response time of the system. For example, both [42] and [43] require image processing and re-classification of the resulting images. Therefore they can be inefficient when dealing with adversarial scenarios such as novelty detection on security camera and image recognition systems on autonomous driving applications which run in real time with high frame rate. Besides, the impact of pre-processing on the classification accuracy is still not fully understood.

Detecting adversarial perturbation indeed can be helpful in practice. However, the fundamental problem is still left unsolved: the neural networks are still not able to recognize similar images as such, ignoring small adversarial perturbation. By proposing novel attack methods, we aim to emphasize the existence of different types of vulnerabilities and the corresponding understanding.

## VI. FUTURE WORK

The DE utilized in this research belongs to a big class of algorithms called evolutionary strategies [48] which includes

other variants such as Adaptive DE [49] and Covariance matrix adaptation evolution strategy (CMA-ES) [50], [51], [52]. In fact, there are a couple of recent developments [61], [62], [63] in evolutionary strategies and related areas that could further improve the current method, allowing for more efficient and accurate attacks.

Furthermore, evolutionary computation also provides some promising approaches to solve adversarial machine learning related vulnerabilities. In fact, evolutionary-based machine learning allows for a great flexibility of models and may be an answer to the same problems it is revealing. First, in an area of evolutionary machine learning called neuroevolution, it was shown to be possible to learn not just the weights but also the topology of the network with evolutionary computation [26], [44], [53]. In fact, SUNA [26] goes beyond current neural models to propose a unified neuron model (e.g., time-scales, neuromodulation, feedback, long-term memory) that can adapt its structure and models to learn completely different problems (including non-markov problems) without changing any of its hyper-parameters. This generality is currently surpassing most if not all deep learning algorithms. Last but not least, self-organizing and novelty-organizing classifiers can adapt to changes in the environment by using flexible representations [54], [55], [56]. For example, they can adapt to mazes that change in shape and to problems where the scope of variables change throughout the experiment [57]: a very challenging scenario in which most if not all deep learning algorithms fail. These among other achievements [58], [59] show a promising path that may solve current problems in deep neural networks in the years to come.

Besides, it can be seen that the one-pixel attack can be potentially extended to other domains such as natural language processing, speech recognition, which will be also left for future work.

## VII. ACKNOWLEDGMENT

This research was partially supported by Collaboration Hubs for International Program (CHIRP) of SICORP, Japan Science and Technology Agency (JST), and Kyushu University Education and Research Center for Mathematical and Data Science Grant.

## REFERENCES

- [1] M. Barreno, B. Nelson, A. D. Joseph, and J. Tygar. The security of machine learning. *Machine Learning*, 81(2):121–148, 2010.
- [2] M. Barreno, B. Nelson, R. Sears, A. D. Joseph, and J. D. Tygar. Can machine learning be secure? In *Proceedings of the 2006 ACM Symposium on Information, computer and communications security*, pages 16–25. ACM, 2006.
- [3] J. Brest, S. Greiner, B. Boskovic, M. Mernik, and V. Zumer. Self-adapting control parameters in differential evolution: A comparative study on numerical benchmark problems. *IEEE transactions on evolutionary computation*, 10(6):646–657, 2006.
- [4] P. Civicioglu and E. Besdok. A conceptual comparison of the cuckoo-search, particle swarm optimization, differential evolution and artificial bee colony algorithms. *Artificial intelligence review*, pages 1–32, 2013.
- [5] H. Dang, Y. Huang, and E.-C. Chang. Evading classifiers by morphing in the dark. 2017.
- [6] S. Das and P. N. Suganthan. Differential evolution: A survey of the state-of-the-art. *IEEE transactions on evolutionary computation*, 15(1):4–31, 2011.
- [7] S. M. Moosavi Dezfouli, F. Alhussein, and F. Pascal. Deepfool: a simple and accurate method to fool deep neural networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2574–2582, 2016.
- [8] S. M. Moosavi Dezfouli, F. Alhussein, F. Omar, F. Pascal, and S. Stefano. Analysis of universal adversarial perturbations. *arXiv preprint arXiv:1705.09554*, 2017.
- [9] A. Fawzi, S. M. Moosavi Dezfouli, and P. Frossard. A geometric perspective on the robustness of deep networks. Technical report, Institute of Electrical and Electronics Engineers, 2017.
- [10] A. Fawzi, S.-M. Moosavi-Dezfouli, P. Frossard, and S. Soatto. Classification regions of deep neural networks. *arXiv preprint arXiv:1705.09552*, 2017.
- [11] I. J. Goodfellow, J. Shlens, and C. Szegedy. Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*, 2014.
- [12] A. Krizhevsky and G. Hinton. Learning multiple layers of features from tiny images. 2009.
- [13] M. Lin, Q. Chen, and S. Yan. Network in network. *arXiv preprint arXiv:1312.4400*, 2013.
- [14] S. M. Moosavi Dezfouli, A. Fawzi, O. Fawzi, and P. Frossard. Universal adversarial perturbations. In *Proceedings of 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, number EPFL-CONF-226156, 2017.
- [15] N. Narodytska and S. Kasiviswanathan. Simple black-box adversarial attacks on deep neural networks. In *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 1310–1318. IEEE, 2017.
- [16] A. Nguyen, J. Yosinski, and J. Clune. Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 427–436, 2015.
- [17] N. Papernot, P. McDaniel, I. Goodfellow, S. Jha, Z. B. Celik, and A. Swami. Practical black-box attacks against machine learning. In *Proceedings of the 2017 ACM on Asia Conference on Computer and Communications Security*, pages 506–519. ACM, 2017.
- [18] N. Papernot, P. McDaniel, S. Jha, M. Fredrikson, Z. B. Celik, and A. Swami. The limitations of deep learning in adversarial settings. In *Security and Privacy (EuroS&P), 2016 IEEE European Symposium on*, pages 372–387. IEEE, 2016.
- [19] A. Rozsa, E. M. Rudd, and T. E. Boulton. Adversarial diversity and hard positive generation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops*, pages 25–32, 2016.
- [20] K. Simonyan, A. Vedaldi, and A. Zisserman. Deep inside convolutional networks: Visualising image classification models and saliency maps. *arXiv preprint arXiv:1312.6034*, 2013.
- [21] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- [22] J. Springenberg, A. Dosovitskiy, T. Brox, and M. Riedmiller. Striving for simplicity: The all convolutional net. In *ICLR (workshop track)*.
- [23] R. Storn and K. Price. Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces. *Journal of global optimization*, 11(4):341–359, 1997.
- [24] S. Christian, Z. Wojciech, S. Ilya, b. Joan, E. Dumitru, G. Ian, F. Rob. Intriguing properties of neural networks. In *In ICLR*. Citeseer, 2014.
- [25] Y. Taigman, M. Yang, M. Ranzato, and L. Wolf. Deepface: Closing the gap to human-level performance in face verification. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1701–1708, 2014.
- [26] D. V. Vargas and J. Murata. Spectrum-diverse neuroevolution with unified neural models. *IEEE transactions on neural networks and learning systems*, 28(8):1759–1773, 2017.
- [27] D. V. Vargas, J. Murata, H. Takano, and A. C. B. Delbem. General subpopulation framework and taming the conflict inside populations. *Evolutionary computation*, 23(1):1–36, 2015.
- [28] D. Wei, B. Zhou, A. Torralba, and W. Freeman. Understanding intra-class knowledge inside cnn. *arXiv preprint arXiv:1507.02379*, 2015.
- [29] J. Yosinski, J. Clune, T. Fuchs, and H. Lipson. Understanding neural networks through deep visualization. In *In ICML Workshop on Deep Learning*. Citeseer.
- [30] M. D. Zeiler and R. Fergus. Visualizing and understanding convolutional networks. In *European conference on computer vision*, pages 818–833. Springer, 2014.
- [31] J. Su, D. Vargas, and K. Sakurai. One pixel attack for fooling deep neural networks. *arXiv preprint arXiv:1710.08864*, 2017.
- [32] N. Hansen. The CMA evolution strategy: a comparing review In *Towards a new evolutionary computation*, pages 75–102. Springer, 2006.
- [33] A. Moustafa, B. Bharathan, S. Mani. Did you hear that? Adversarial Examples Against Automatic Speech Recognition. *arXiv preprint arXiv:1801.00554*, 2018.
- [34] P. Nicolas, M. Patrick, S. Ananthram, H. Richard. Crafting Adversarial Input Sequences for Recurrent Neural Networks. *arXiv preprint arXiv:1604.08275*, 2016.
- [35] G. Kathrin, P. Nicolas, M. Praveen, B. Michael, M. Patrick. Adversarial Perturbations Against Deep Neural Networks for Malware Classification. *arXiv preprint arXiv:1606.04435*, 2016.
- [36] E. Logan, T. Brandon, T. Dimitris, S. Ludwig, M. Aleksander. A Rotation and a Translation Suffice: Fooling CNNs with Simple Transformations. *arXiv preprint arXiv:1712.02779*, 2017.
- [37] G. Kathrin, P. Nicolas, M. Praveen, B. Michael, M. Patrick. Adversarial Perturbations Against Deep Neural Networks for Malware Classification. *arXiv preprint arXiv:1606.04435*, 2016.
- [38] M. Sharif, L. Bauer, MK. Reiter. On the Suitability of Lp-norms for Creating and Preventing Adversarial Examples. *arXiv preprint arXiv:1802.09653*, 2018.
- [39] X. Yuan, P. He, Q. Zhu, R. R. Bhat. Adversarial Examples: Attacks and Defenses for Deep Learning. *arXiv preprint arXiv:1712.07107*, 2017.
- [40] N. Papernot, P. McDaniel, X. Wu, S. Jha, A. Swami. Distillation as a defense to adversarial perturbations against deep neural networks. In *Proceedings of IEEE Symposium on Security and Privacy (SP)*, pp 17011708.
- [41] R. Huang, B. Xu, D. Schuurmans, C. Szepesvri. Learning with a strong adversary. *arXiv preprint arXiv:1511.03034*, 2015.
- [42] W. Xu, D. Evans, Y. Qi. Feature squeezing: Detecting adversarial examples in deep neural networks. *arXiv preprint arXiv:1704.01155*, 2017.
- [43] B. Liang et al. Detecting Adversarial Examples in Deep Networks with Adaptive Noise Reduction. *arXiv preprint arXiv:1705.08378*, 2017.
- [44] K. O. Stanley, R. Miikkulainen. Evolving neural networks through augmenting topologies. In *Evolutionary Computation*, pp. 99–127.
- [45] N. Carlini, D. Wagner. Adversarial examples are not easily detected: Bypassing ten detection methods. In *Proceedings of the 10th ACM Workshop on Artificial Intelligence and Security*, pp. 3–14.
- [46] N. Carlini, D. Wagner. Defensive distillation is not robust to adversarial examples. *arXiv preprint arXiv:1607.04311*, 2016.
- [47] N. Carlini, D. Wagner. Towards evaluating the robustness of neural networks. In *2017 IEEE Symposium on Security and Privacy (SP)*, pp. 39–57.
- [48] H.G. Beyer and H.P. Schwefel. Evolution strategies: A comprehensive introduction. In *Natural computing*, pp.3–52.
- [49] A.K. Qin and P.N. Suganthan. Self-adaptive differential evolution algorithm for numerical optimization. In *The 2005 IEEE Congress on Evolutionary Computation*, pp. 1785–1791.
- [50] N. Hansen, S. D. Miller, and P. Koumoutsakos. Reducing the time complexity of the derandomized evolution strategy with covariance matrix adaptation (CMA-ES). In *Evol. Comput.*, pp. 1–18.
- [51] N. Hansen and A. Ostermeier. Adapting arbitrary normal mutation distributions in evolution strategies: The covariance matrix adaptation. In *Proceedings of the 1996 IEEE Conference on Evolutionary Computation*, pp. 312–317.
- [52] N. Hansen and A. Ostermeier. Completely derandomized self-adaptation in evolution strategies. In *Evol. Comput.*, pp.159–195.
- [53] D. Whitley, S. Dominic, R. Das and C.W. Anderson. Genetic reinforcement learning for neurocontrol problems. In *Machine Learning*, pp.259–284.
- [54] D.V. Vargas, H. Takano and J. Murata. Self organizing classifiers: first steps in structured evolutionary machine learning. In *Evolutionary Intelligence*, pp.57–72.
- [55] D.V. Vargas, H. Takano and J. Murata. Self organizing classifiers and

- niched fitness. In *Proceedings of the 15th annual conference on Genetic and evolutionary computation*, pp. 1109-1116.
- [56] D.V. Vargas, H. Takano and J. Murata. Novelty-organizing team of classifiers-a team-individual multi-objective approach to reinforcement learning. In *Proceedings of the SICE Annual Conference (SICE)*, pp. 1785-1792.
  - [57] D.V. Vargas, H. Takano and J. Murata. Novelty-organizing team of classifiers in noisy and dynamic environments. In *2015 IEEE Congress on Evolutionary Computation (CEC)*, pp. 2937-2944.
  - [58] R.J. Urbanowicz and J.H. Moore. ExSTraCS 2.0: description and evaluation of a scalable learning classifier system. In *Evolutionary intelligence*, pp.89-116.
  - [59] I.M. Alvarez, W.N. Browne and M. Zhang. Compaction for code fragment based learning classifier systems. In *Australasian Conference on Artificial Life and Computational Intelligence*, pp. 41-53.
  - [60] N. Carlini and D. Wagner. Magnet and efficient defenses against adversarial attacks are not robust to adversarial examples. *arXiv preprint arXiv:1711.08478*, 2017.
  - [61] A. Abdolmaleki, B. Price, N. Lau, L.P. Reis and G. Neumann. Deriving and improving CMA-ES with information geometric trust regions. In *Proceedings of the Genetic and Evolutionary Computation Conference*, pp. 657-664.
  - [62] K. Nishida and Y. Akimoto. PSA-CMA-ES: CMA-ES with population size adaptation. In *Proceedings of the Genetic and Evolutionary Computation Conference*, pp. 865-872.
  - [63] M. Groves and J. Branke. Sequential sampling for noisy optimisation with CMA-ES. In *Proceedings of the Genetic and Evolutionary Computation Conference*, pp. 1023-1030.